

Final report for the work done in lupulo

Alejandro López Espinosa

January 20, 2016

Abstract

I developed a framework to build real time web pages as my final project. In this document I explain how the project was born and developed through the five months of internship at the VIVES university.

Contents

1	Index of words	1
2	Introduction	2
2.1	Context	2
2.2	General idea	2
3	Analysis	3
3.1	Goals	3
3.2	Requirements	3
3.2.1	Functional requirements	3
3.2.2	Non functional requirements	5
3.2.3	Usability requirements	6
3.3	Users	7
3.4	Technologies	7
4	Design	8
4.1	Data schema language	9
4.2	Layout language	9
4.2.1	Inheritance	10
4.3	Templates	11
4.4	Widgets	12
4.5	Accessors	12
4.6	Listeners	13
4.7	Workflow	13
5	Architecture	14
5.1	Static structure	14
5.2	Dynamic structure	17
6	Methodology of work	22
7	Work to do	22
8	Conclusions	23

1 Index of words

Accessor: Javascript abstraction that allows a widget to access the data sent by the device and defined in the data schema definition without knowing that definition.

Backend: The server side run by the programmer of the web page which pushes information in an asynchronous way to the frontend.

Frontend: The real time web page built by the framework from the descriptions given by the programmer of the web page.

Jinja2: Python module that provides a synchronous template system.

Layout: Description of the widgets in the web page.

Listener: Python object that is listening to the information coming from the device.

Lupulo: Name of the final project.

Pip: Python Package Index, the central repository of modules for the python environment.

Server Sent Events: Html5 API that provides a way to push data to a web page.

SSE: Acronym for Server Sent Events.

Template: Jinja2 template used to render an html web page.

Twisted: Asynchronous networking framework built in python2.

Widget: Javascript object that renders some information coming from the backend.

WSGI: The Web Server Gateway Interface (WSGI) is a specification for simple and universal interface between web servers and web applications or frameworks for the Python programming language.

2 Introduction

2.1 Context

The first proposal for my final project was about building a web page that could monitor and command a pololu robot.

Once I started building that web page I realised that a lot of the code I was writing was very generic and that the main mechanism to push information to the web page was the same not only for that robot in particular but for any device that needs to be monitored in real time.

So, I started thinking about a framework to build real time web pages where the focus was on providing a way to describe in a very easy way both the data the device is sending and how that data should be visualized in the web page. That is the main goal that my final project achieves.

2.2 General idea

There are several principal use cases that the framework should provide support for, but the main one was to allow a programmer to set up a real time web page in the easiest and most personalized way possible.

I achieved the easiness with a couple of high level languages that define the dynamic data that the device is sending and the dynamic way of visualizing that data in the web page. Both languages allow an easy construction of a real time web page focused on visualizing information.

I achieved the customization of the web page by providing mechanisms to program the web page with the usual tools like html, javascript or css. For even further customization, lupulo also provides several mechanisms to extend its behaviour with custom python code that the framework calls if it's present in the project directory and that changes the way the framework works.

The general philosophy behind lupulo is to allow people with very different levels of experience to build real time web pages. The unexperienced users can use all the default mechanisms that lupulo provides to describe in a high level a web page. In the same manner, experienced users can use their extra skills to personalize the way lupulo works and build more complex real time web pages while at the same time using the high level descriptions to abstract away the general mechanisms that make their web page different to others.

In the following chapters I'm going to describe in plain English how the analysis

of the problem drove the design of the solution and its implementation in a concrete architecture with a set of defined technologies.

3 Analysis

The first stage of any project is the analysis of the problem that the software is going to solve. In the following sections I focus on the problem lupulo solves describing it's basic components.

3.1 Goals

As already described above, the two most important goals are to provide an easy and customizable tool to build real time web pages.

For the tool to be used it's important to allow programmers to build very fast prototypes of the final web page maybe only to test how the device they are monitoring behaves. It's also important to provide an easy way to record all the information the device is sending because that information is very valuable.

In the same way, it's very important that the software grants a bunch of customizations that allow the programmer to build much more complex web pages where the consumed time for building them is not so important but the final web page is.

Sometimes the two goals collide and it's impossible to provide an easy and also extensible tool. A lot of that situations has been studied in this document and a solution has been designed and implemented in lupulo.

3.2 Requirements

Although the usability requirements are also non functional, I have decided to separate them due to their amount and importance.

3.2.1 Functional requirements

Title	Asynchronous data link between the frontend and the backend.
Description	The backend must push the information it receives from the device as long as it is available.

Title	Widgets to visualize information.
Description	The frontend must provide an abstraction that allows the easy description of a real time web page for visualization of information.
Title	High level description of the data schema.
Description	The data schema that the device is sending must be described dynamically.
Title	Infinite recursion in the data schema language.
Description	A descriptor in the data schema can be composed of infinite definitions of other descriptors.
Title	High level description of the widgets in the web page.
Description	The visualization widgets of the web page must be described entirely in the backend.
Title	Independence of the data schema.
Description	The data the device is sending to the framework can change very often so a dynamic description of the data schema should be provided.
Title	Independence of the data link with the device.
Description	Different data links between the backend and the device should be allowed.
Title	Hot layout and data schema definition.
Description	When the data schema or layout of the project change, the web page should change accordingly.
Title	Templates mechanism.
Description	The backend should provide a way to write web pages with templates.
Title	Static files from lupulo and project directory.
Description	Both the lupulo framework and the specific project built upon it should be able to serve static files.
Title	Custom urls with support to RESTful interfaces.
Description	The backend will provide a custom RESTful interface that the user can change.
Title	Dynamic controllers in the frontend.
Description	The main controller of the web page must provide a mechanism to be redefined or overwritten completely.

Title	Alerts mechanism in the frontend.
Description	The frontend should provide a way to post alerts to the web page.
Title	Independence between the widget and the data schema definition.
Description	The widget cannot know the schema definition for one specific project, so the frontend must provide a way for the widget to access the data in its paint method.
Title	Debug web page.
Description	A debugging web page should be provided that allow the programmer to easily track down an error in the widget, the data schema or the layout.
Title	SSE standalone client.
Description	An standalone client should be provided that allow the programmer to see directly the data that the backend sends to the frontend.
Title	Responsive web page.
Description	The real time web page should be responsive to different size factors.

3.2.2 Non functional requirements

Title	Multiple parallel http requests.
Description	The backend must be able to serve multiple http requests.
Title	Multiple devices.
Description	The backend must be able to track several devices connected at the same time.
Title	The server must be run under a RaspberryPi.
Description	The resources of a RaspberryPi must be enough to run the server.
Title	The rendering of a template cannot block the sse channel.
Description	If a template is being rendered for one web page, it cannot block the pushing mechanism of another http clients.
Title	The storing of the information in a persistent storage cannot block a http request.
Description	The backend must store the data without interfering with the web server.

Title	The server shouldn't be executed as a superuser.
Description	In order to allow a sensible use of the program, the process should not require superuser permissions.
Title	The delay since some data is pushed into the frontend until it's rendered is under the second.
Description	In order to grant the real time rendering of information, the admissible delay of rendering some information should be around 1 second.
Title	The delay since some data is sent by the device until it's received by the framework is under 1s.
Description	In order to grant the real time delivery of information, the admissible delay of delivering some information should be around 1 second.

3.2.3 Usability requirements

Title	Creation of a web page under 1 hour.
Description	A user who haven't built any web page with lupulo but who is an experienced web developer should be able to build a functional web page in less than one hour.
Title	Refresh of the web page without restarting the server.
Description	In order to allow a smooth workflow, it shouldn't be necessary to restart the web server in order to see the changes made to a template.
Title	Base templates and style sheets.
Description	The framework should provide a bunch of base templates and style sheets that the project can be built upon.
Title	Defaults files when the project is created.
Description	To allow beginners to understand better the framework, some default files should be placed when the project is started.
Title	Provide several widgets with the installation of the framework.
Description	In order to build complex visualizations very fast, the widgets for that visualizations should be provided with the basic install.

Title	Global commands for creation of the project and for launching of the server.
Description	At least two commands should be provided to create the project and to launch a web server from that project directory.

3.3 Users

There are two major type of users of the system. The first one is the user of the framework, which is responsible of writing the data schema and layout descriptions, also of launching the server and expanding the framework if necessary.

This user can have several levels of skill but I'll focus on two different users:

- the beginner user who want to use the framework to visualize and record the data the device is sending as soon as possible.
- the experienced user who want to build a complex real time web site.

Also, the final user of the real time web page is a user of the framework because it also has an impact on the development of the framework.

3.4 Technologies

In order to build a web server that could push information to the frontend in an asynchronous way, several technologies were studied, mainly web sockets and the Server Sent Event API.

Both of them are html5 APIs. The main difference between the two is that web sockets provides a bidirectional channel between the web page and the server whereas SSE only provides a unidirectional channel to push information from the server to the web page.

After a little thought, SSE was chosen because it's more easy to work with it in javascript and because it provides by default a built-in way of demultiplexing several channels of information that share the same SSE connection, a feature I needed.

Once I decided to use SSE, I needed to find a framework that would allow me to build an asynchronous web server that used SSE to push information to the web page.

I also pay some attention to the fact that maybe in the future I needed to switch from SSE to web sockets because the first technology was not good enough, so I decided to choose a framework that provided support for both technologies.

I was very familiarized with Twisted and it provided built-in support for both technologies so it was a natural decision for me.

Once I decided to use twisted as the framework, I have to decide between using a WSGI application as the main engine for the web server or build my own asynchronous engine. After digging a little more into how WSGI applications are integrated in twisted, I discovered that the http requests wouldn't be served asynchronously if I added a WSGI application to the twisted application, so I decided to use my own web server instead of using one of the existing WSGI frameworks in the community to provide more efficiency.

Another important decision concerning technologies is the use of jinja2 as a templates system for the web server due to the fact that I wasn't using a WSGI application. In order to allow multiple http requests and devices, it was necessary to render the web pages in an asynchronous way to avoid the blocking of requests, so I built around jinja2 a system to asynchronously render a web page.

4 Design

As already explained, lupulo is supposed to allow programmers to build real time web pages in an easy and customizable way.

The design of lupulo is focused around several abstractions that allow the user to build a real time web page. These abstractions are split between beginner and expert abstractions so that the beginner ones provide the easiness of use and the expert ones the required customization.

The beginner abstractions are the layout language, the data schema language and the templates. The expert abstractions are the widgets construction, the accessors and the listeners.

All the abstractions are supposed to work together in a smooth workflow that the user of the framework can follow in order to build the real time web pages.

In the following sections I'll explain every abstraction in isolation first and then how everything works together to provide a smooth workflow.

This is a short description of the abstractions, if you want to know more you should read the [official documentation of the project](#).

4.1 Data schema language

The data schema provides a dynamic way of describing the data the device is sending that allows the backend to process the data and automatically verify and generate (for debugging purposes) it. This abstraction also supports the accessor mechanism described below.

The data schema language is a JSON file that describes the schema of the data received by the backend. This file must have a global JS object that will be interpreted as the data schema. This global object has several sections called data sources.

A data source is a stream of data that a widget in the frontend can subscribe to. It's usually paired with the concept of sensor. The data measured by one sensor is usually packaged into one data source that the widgets can listen to, but you can design weird or more complex data schemas. So it's up to you how the data schema looks like as long as it's valid.

Each widget in the frontend will listen to one or more data sources from the device and will construct the visualizations with that information.

Each data source is described by a JSON object (in the data schema file) filled with some key-value pairs that describe the properties of the data associated with a data source. There is only one obligatory property of every data source: its type.

The type of a data source determines how the data is going to be verified by the backend when it is going to be forwarded to the frontend plus some other less interesting things. Each type defines its own properties that every data source description of that type must stick to.

Currently there are two kind of types:

- Primitive types
- Aggregated types

The aggregated types are a composition of primitive types and the primitive types describe an atom of data like a number, a date or an enumerated value.

4.2 Layout language

A layout is an object that the frontend passes to a widget when it's going to be constructed. The main responsibility of a layout is to let the widget know what and how to render the information that the device is sending him. Also a layout provides a description to define what widgets to construct or where they should be placed

into the web page.

The layout file is a JSON file which lists all the layouts that are needed to render a web page. A layout is filled with data that:

- The backend and frontend need to construct the appropriate widget for a given layout.
- The widget needs itself to render the information.

Every layout thus shares some common attributes that they must provide in order to be valid lupulo layouts. But depending on the type of widget that the layout is describing, another set of attributes are needed. These last attributes are defined by the widget.

One key attribute is the `event_names` attribute, that describes the data sources the widget is listening to. It is a list containing the names of the data sources the widget is listening to.

Another very important attribute is the `type` attribute, that defines which type of widget will be defined by the layout.

Other obligatory attribute is `size` which is a dictionary with two keys: `height` and `width`. The size defines, as his name implies, the size of the widget in the web page.

In order to add the widget to the web page, the web page needs to have a html element to bind the widget to. This element is called the anchor of the widget in the web page and is another obligatory attribute of every layout in the layout file. This anchor can be any selector of jquery that resolves to an html element.

As described above, apart from this attributes that every layout must define, each type of widget defines its own set of attributes, so a layout must also provide this set of attributes in order for the layout to be valid.

4.2.1 Inheritance

Due to the verbosity and duplicity of the attributes of a layout, the layout language provides a way to inherit attributes from a parent. Every attribute defined in a parent that isn't defined in a child will be inherited in the child. If an attribute exists both in the child and parent, the child value for the attribute overwrites that of the parent in the child definition.

There are two kinds of layouts:

- Leaf layout: a layout that will be sent to the frontend to construct a widget.
- Abstract layout: a layout that won't be sent to the frontend but that can be

used inside the layout file.

A parent must be abstract in order to be considered for inheritance.

So, a leaf layout will be sent to the frontend to construct a given widget that is described by the direct attributes of his layout and that ones of his parent. However, the abstract layout won't be sent to the frontend so no widget will be constructed with the possibly partial attributes of the parent layout.

To describe that a layout is abstract, you must add the abstract boolean attribute to the layout with a value of true.

To describe that a layout inherits from another one, you must add the parent string attribute to the layout with the value of the name attribute of the parent.

It isn't allowed multiple inheritance but several levels of inheritance are possible.

4.3 Templates

You can also write your own web pages and serve them in a custom sitemap using the templates abstraction.

A template is a jinja2 text file in the templates directory of the project, which will be compiled by the backend to html and served to the user in a specific url that you define in the urls.py file of your project.

The urls.py file must contain a list named urlpatterns made of tuples of two elements. The first element is a string that defines the url that a **twisted resource** (which is the second element) will listen to.

This resource must inherit from `lupulo.http.LupuloResource` and can ask for some template with the method `get_template`, which will return a valid template with the render method.

Once the resource has the template, it can call template's render method with a dictionary context to render some custom information and return what this method returns when it's called by the backend for an http request.

A user's template can inherit from a lupulo template and define only the information relevant to a project in particular. These lupulo templates define a controller that will take care of the construction and destruction of all the widgets defined in the layout file. A user can also overwrite some of the functionality of this controllers to provide a more customized management of the widgets.

4.4 Widgets

You can extend the way the information is displayed in the web page by writing custom widgets. A widget is a javascript object that is constructed by the frontend with a layout description written by the user and sent by the backend. Every widget is connected to some data source in order to create a visualization of some information. That connection is specified in the layout of the widget as already explained above.

Currently, in order to paint something to the screen, the widget can only use the d3.js library.

The lifetime of a widget starts when the frontend receives a valid layout from the backend and the widget is constructed by creating a new object of that type.

Then, the constructor of the widget must call its supertype with `Widget.call(this, layout)`. After this call, the widget will have a bunch of attributes that allow it to interact with the web page in order to render some information in it.

In order to render information, every second the widget's `paint` method will get called with an object filled with the new data pushed by the backend if it's available when the `paint` method is called. In this method, the widget should update its drawings with the new information and the attributes to access the web page.

Furthermore, when the frontend wants to restart a widget, it'll call its `clear_framebuffers` method. This method must be implemented by the widget and must clear any drawing the widget did in the web page.

A widget can aggregate another widgets to render some information by building them manually and then calling its `paint` and `clear_framebuffers` methods whenever it is appropriate.

4.5 Accessors

In order for the widget to access the data received from the device, the widget must know the data schema of the data. But usually the programmer of a widget is a different person than the programmer of a web page, so the data schema is not usually the same among different projects and therefore the widget cannot access directly the data because it doesn't know its schema.

The accessors abstraction allow a widget to access the data without knowing its schema. The widget delegates the real access of the data to another object that is constructed with the data schema and that retrieves the data the widget needs.

The idea is to allow the programmer of the web page to describe in the layout

the data it wants a widget to render paying attention to the widget needs and the structure of the data.

The widget will construct an accessor object with that definition when it's constructed itself. Then, when its paint method is called, it will call the accessor object with the raw data expecting the structured data to be returned as a result. Finally, the widget will draw something in the web page with that structured data.

4.6 Listeners

In order to connect the backend to the data source, the user can build its own listener that will retransmit to the backend the data it receives. A listener is a twisted service that will be run when the server is started.

Once the listener is created by the backend, it will publish data it receives from the device to a twisted resource that will validate and push the information to the frontend.

Therefore, a programmer can customize the way the device is communicating with the exterior world just by writing the correct listener.

4.7 Workflow

In order to build a real time web page with lupulo, a beginner user only has to define a data schema that describes the data its device is sending, a layout of the widgets it want to see at the web page and a settings file that defines some of the behaviour of the server.

With that three basic files, it will have a real time visualization of the data the device is measuring from its surroundings.

If the programmer wants to provide a more customized web page, it can use the templates system to add information to the web pages and customize them also with custom style sheets and javascript source files.

Furthermore, the user can define its own widgets, register them against lupulo and use them in the layout file to get a much more customized web page.

Finally, if the device has a very specific way of communicating to the exterior, the programmer can always define a custom listener that will push information from the device to the backend in a custom way.

In order to ease the debugging of the layout file, the user can use a debug web page that will render in real time the layout file and the data the frontend receives

from the backend in raw form and all of the widgets the frontend construct from the layout definition.

Also, to ease the debugging of listeners or of the layout in more low level, the user of the framework can use a standalone sse client that will print to the standard output the raw information it receives from the backend.

So the basic workflow is to define a layout and a data schema files, run the server and see how everything looks both in the real and debug web pages, then modify the layout or data schema and iterate in this process as much as wanted.

If the programmer needs more customization of the web page, he can write a new widget and use it in the layout file or write a template that renders some static information, or if he wants to communicate to the device in some specific way, he can extend the listeners provided by lupulo. The point is that he can use more advanced abstractions to customize his project.

5 Architecture

Once the basic design of the system has been explained, it's moment to describe the major software components of the system and how they interact with each other.

Before that explanation though, I'll overview the two main technologies that make all of this project possible.

5.1 Static structure

The project is divided between the backend and frontend. The backend is made of the following components:

- A twisted web application that launches the web server and listener.
- The manager of the layout file.
- The manager of the data schema file.
- The manager of the listeners.
- An asynchronous template system.

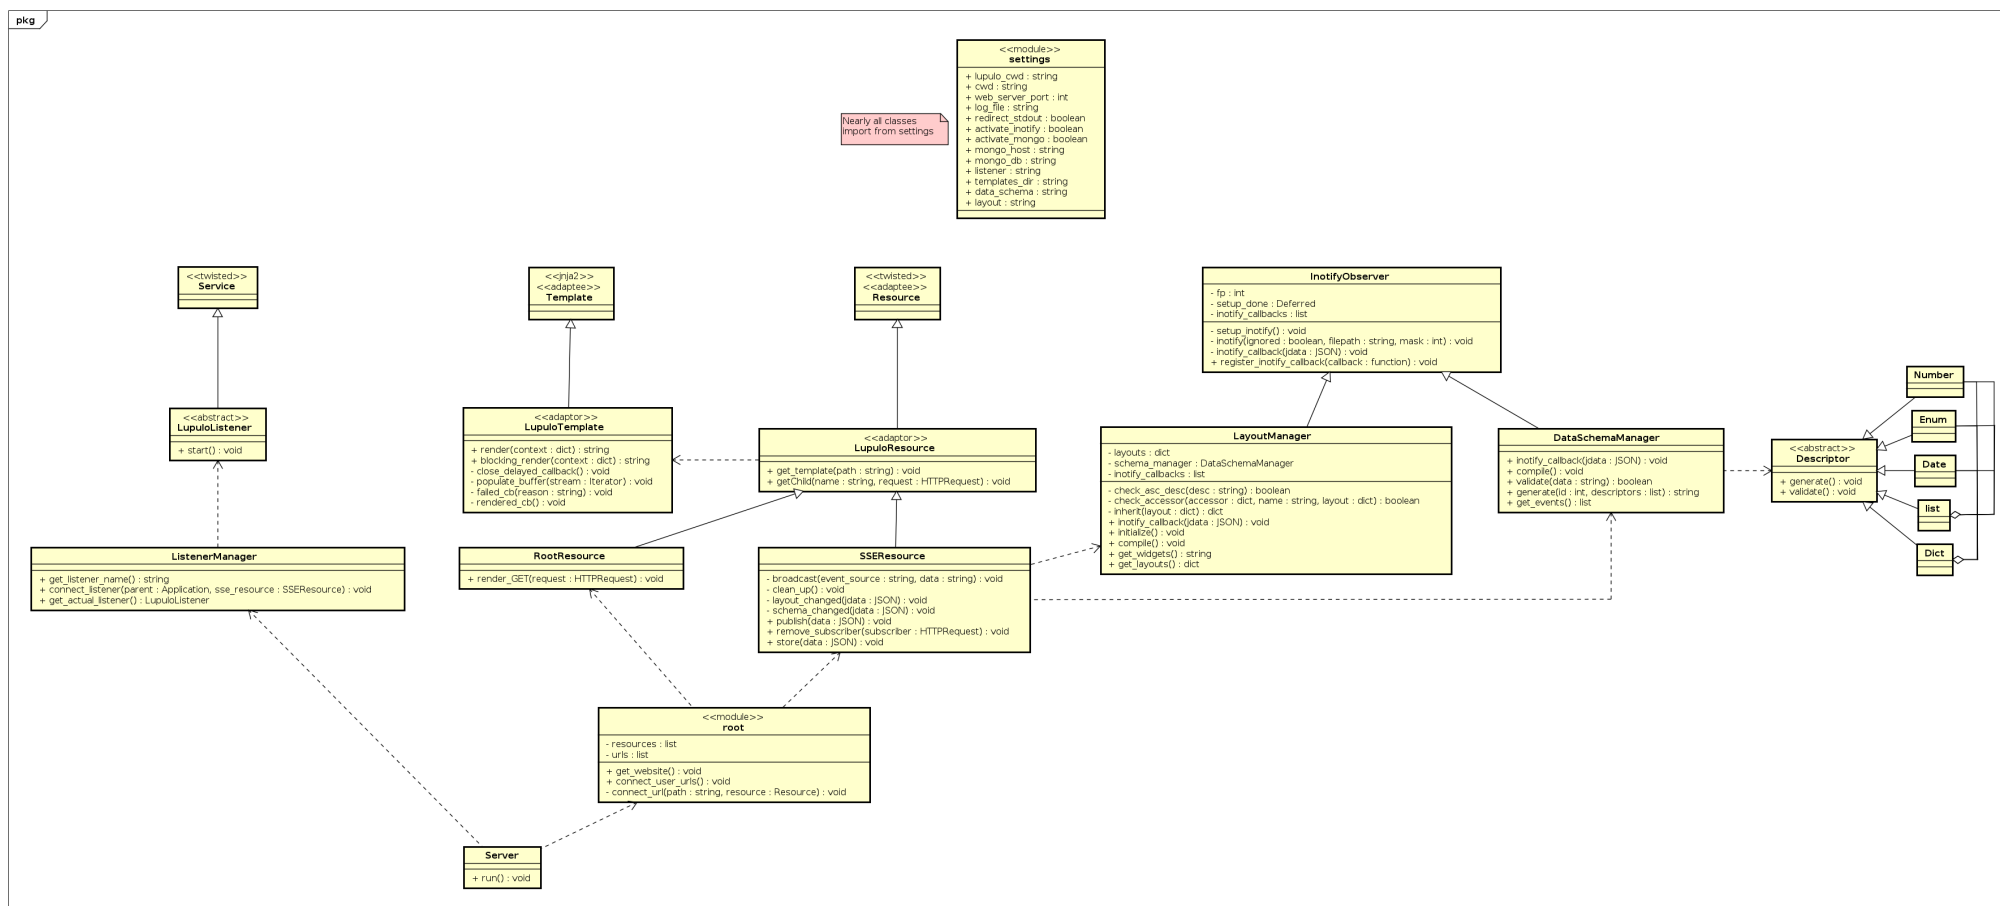


Figure 1: Classes diagram for the backend

The frontend is made of the following components:

- Controller for the events sent by the backend.
- Widget that renders information sent from the backend.
- Accessors abstraction.

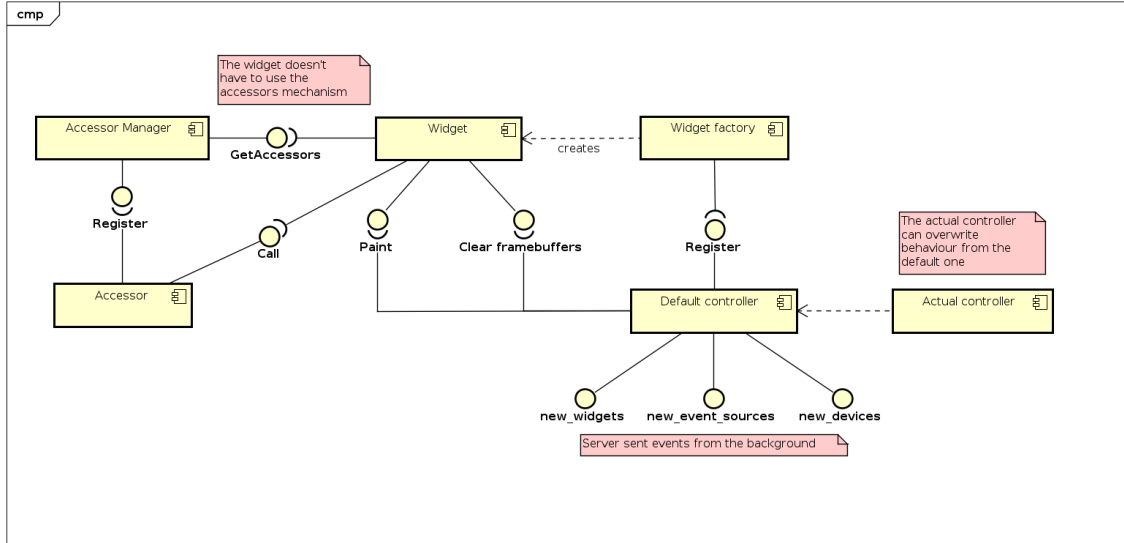


Figure 2: Main components for the frontend

The usual deployment of the project involves several hosts which run the backend, the optional store mechanism and the browser with the frontend loaded on it as can be seen in the following diagram.

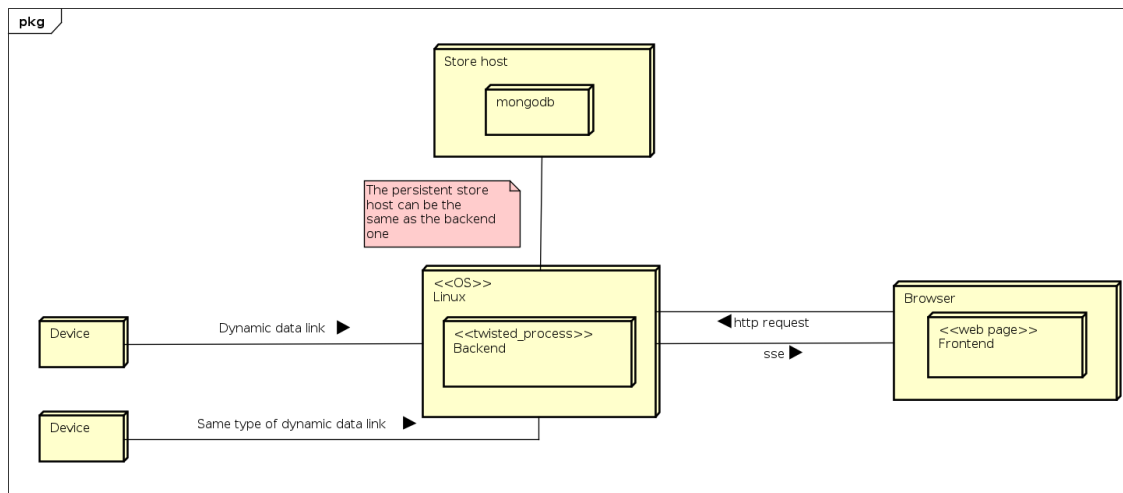


Figure 3: General deployment diagram

5.2 Dynamic structure

Once the main structure of the project has been described, I'll explain the interactions of the software with several sequence diagrams that cover the main stages of lupulo once it has been started.

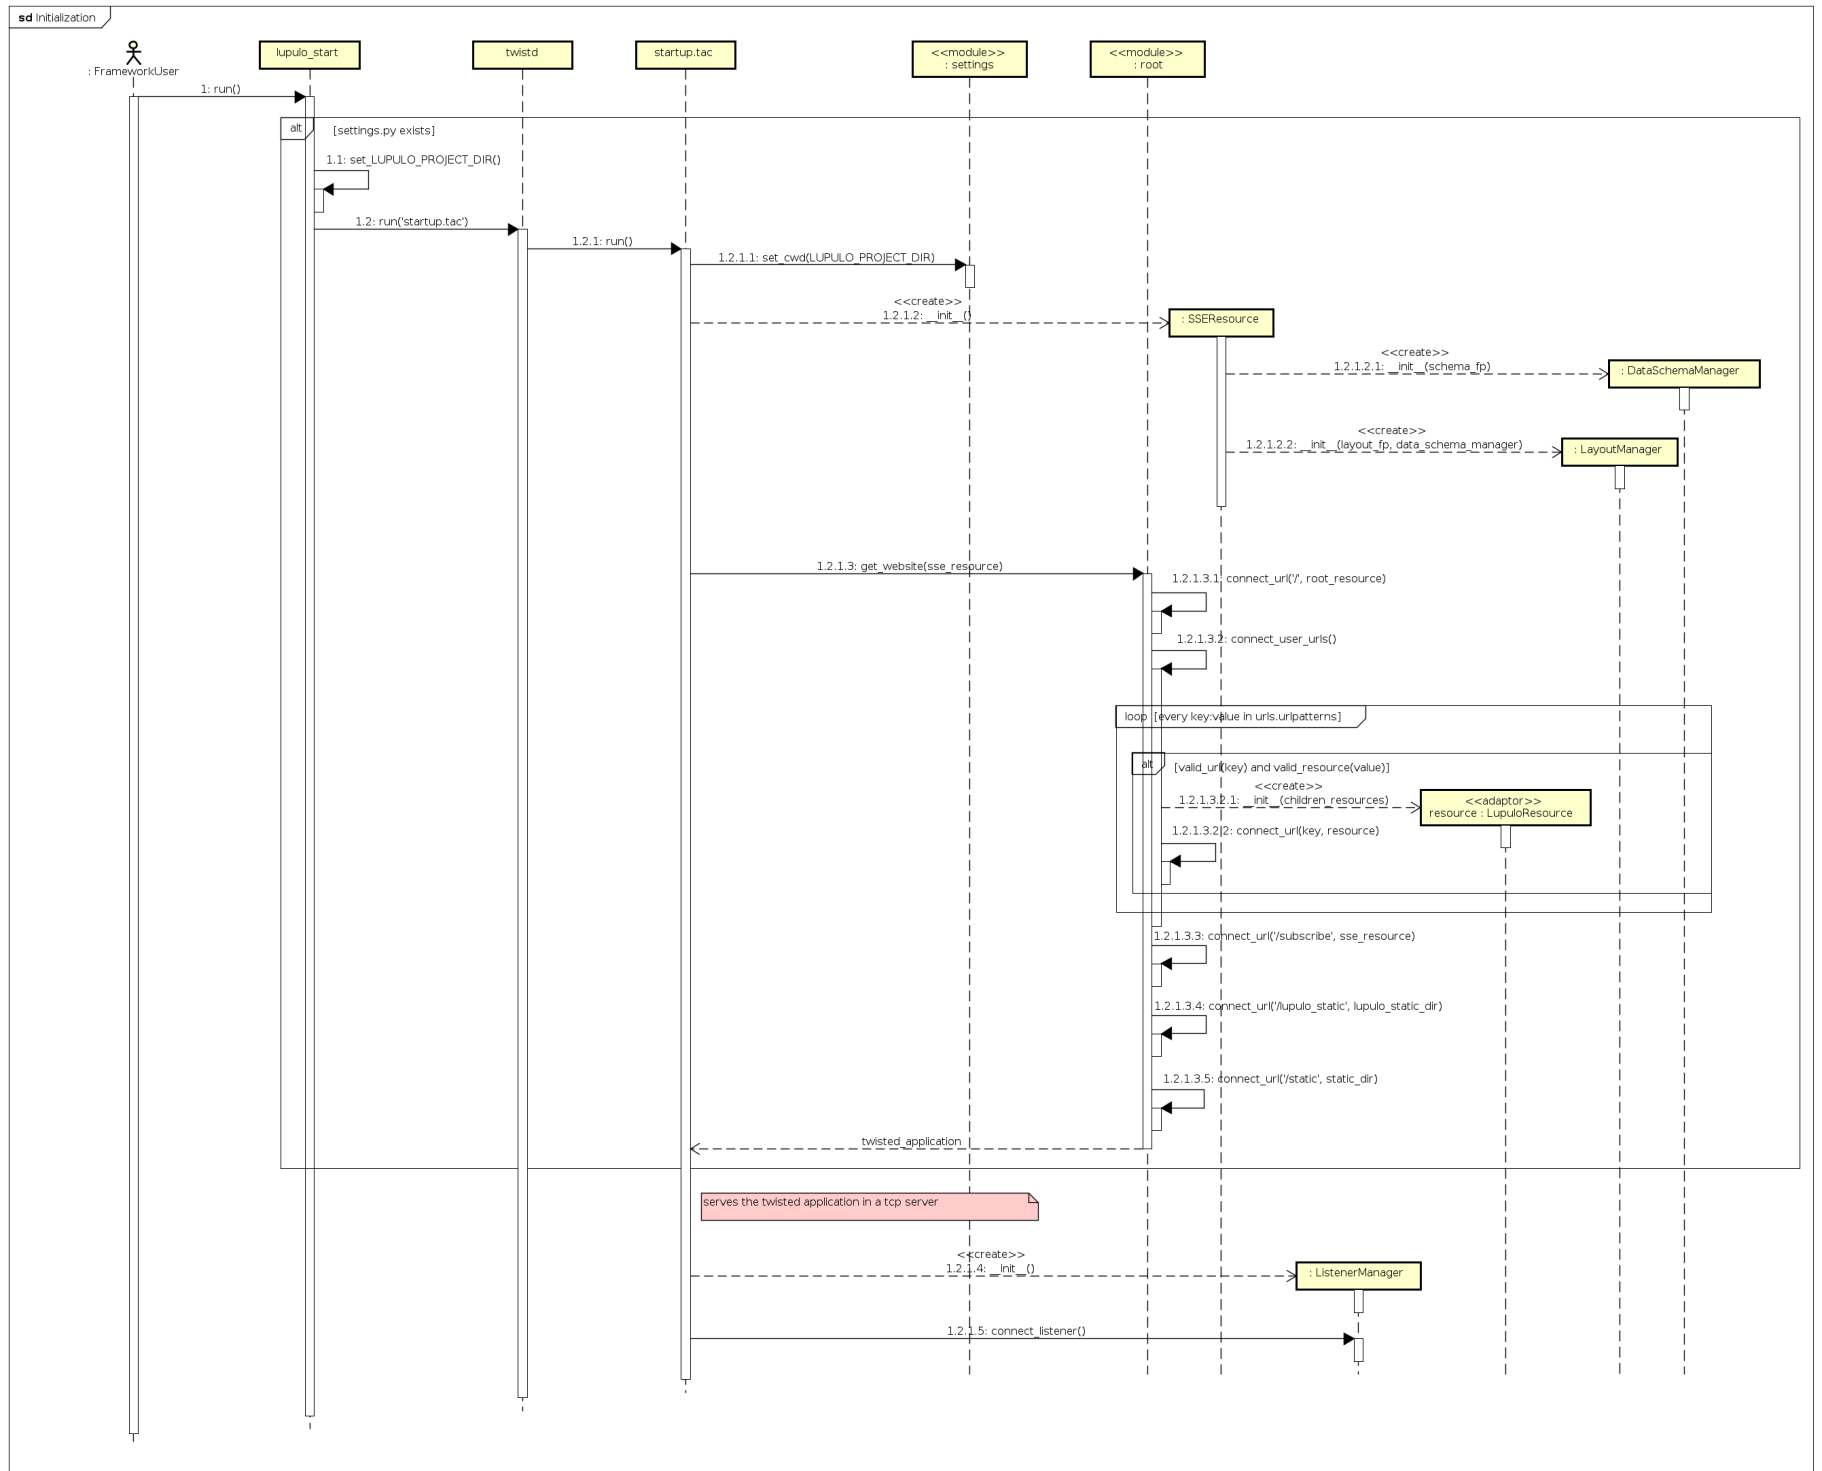


Figure 4: Initialization of the backend

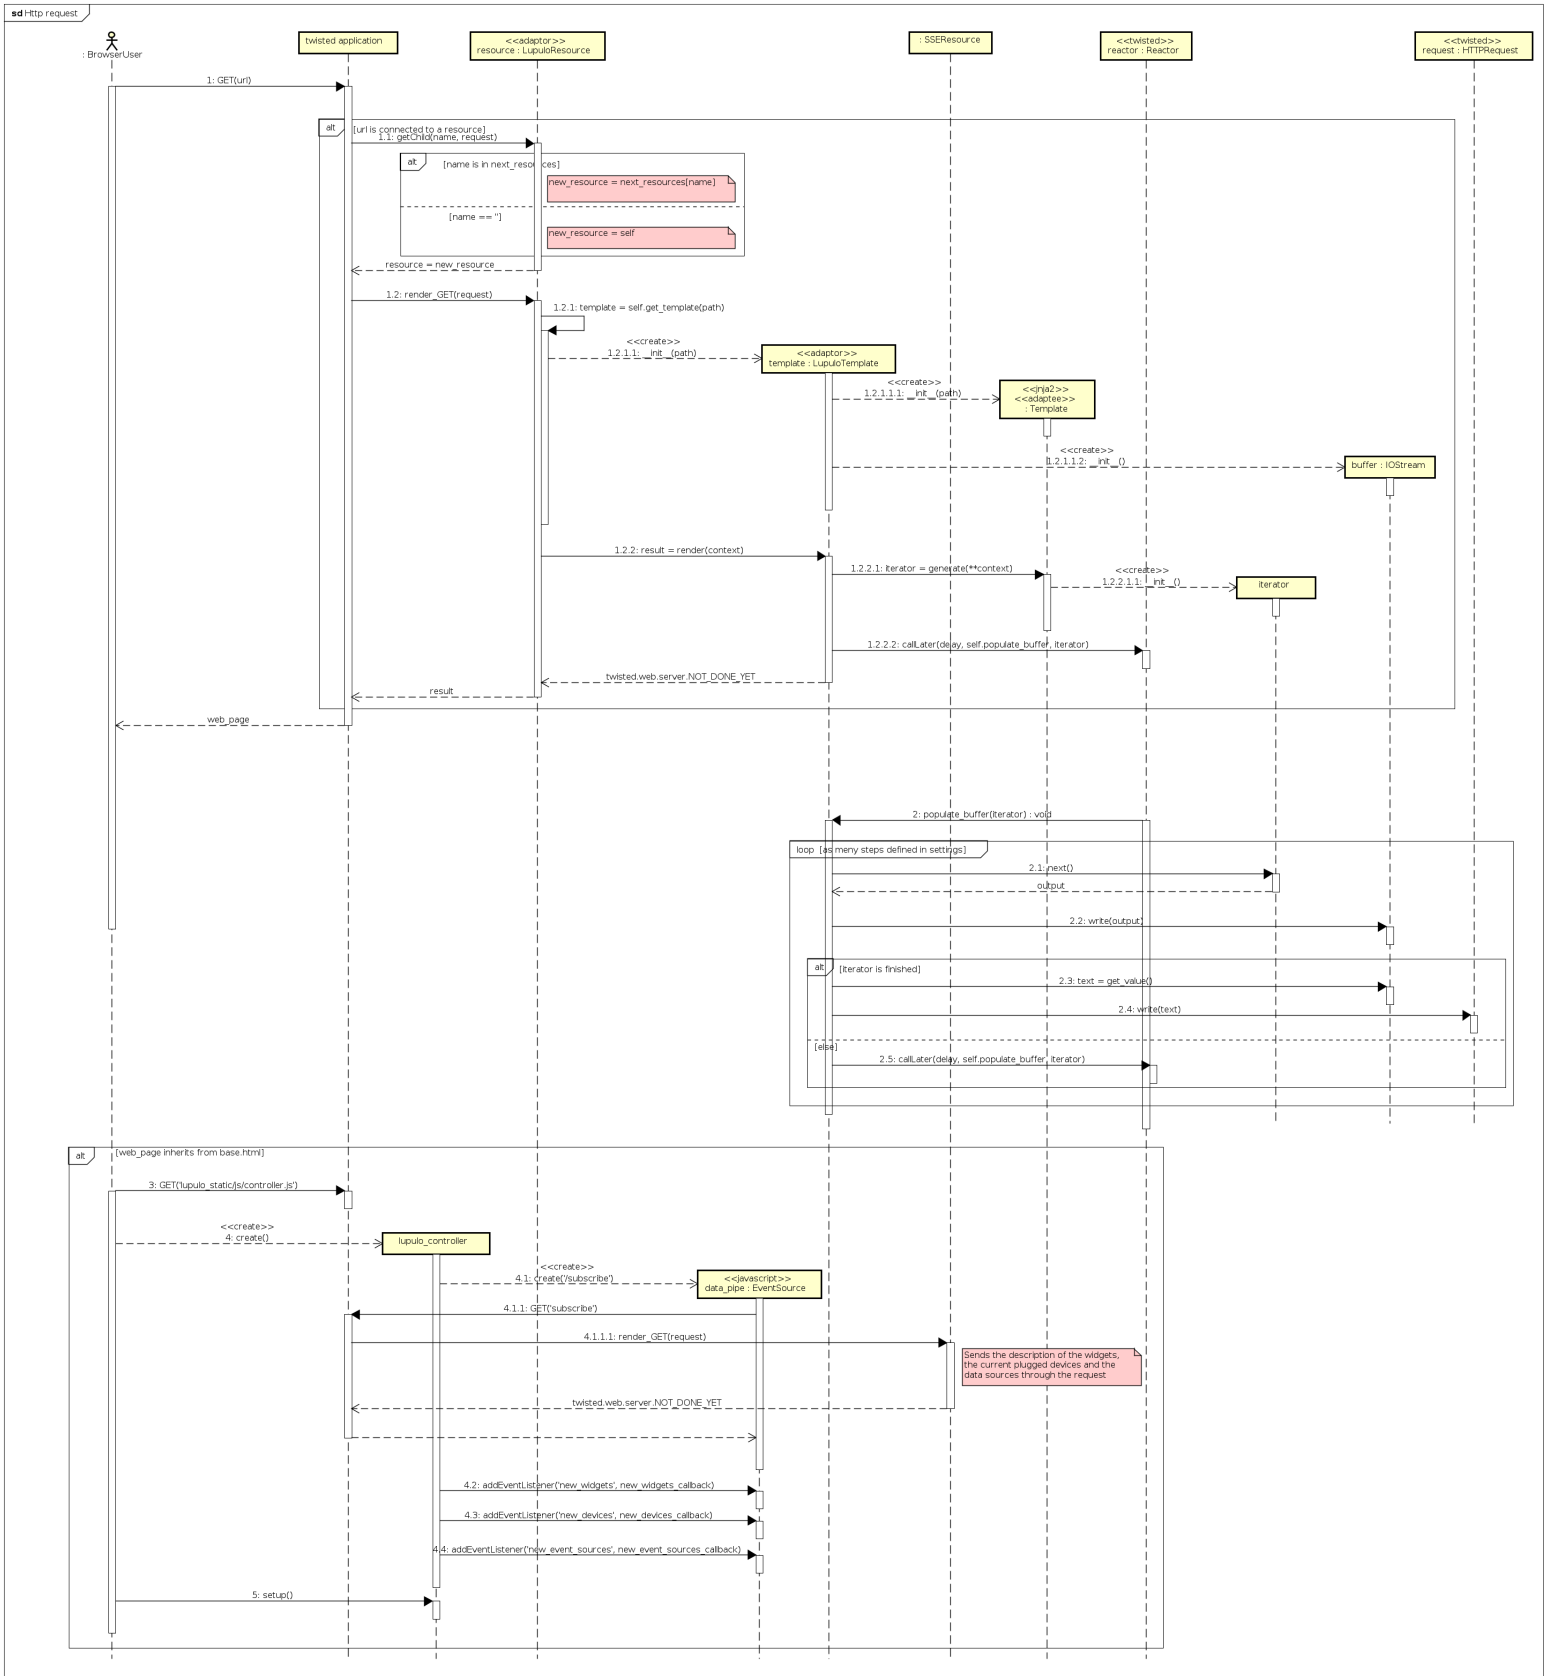


Figure 5: Http request from a browser asking for a web page derived from base.html




Figure 6: Reception of a backend sse event, in particular new_widgets

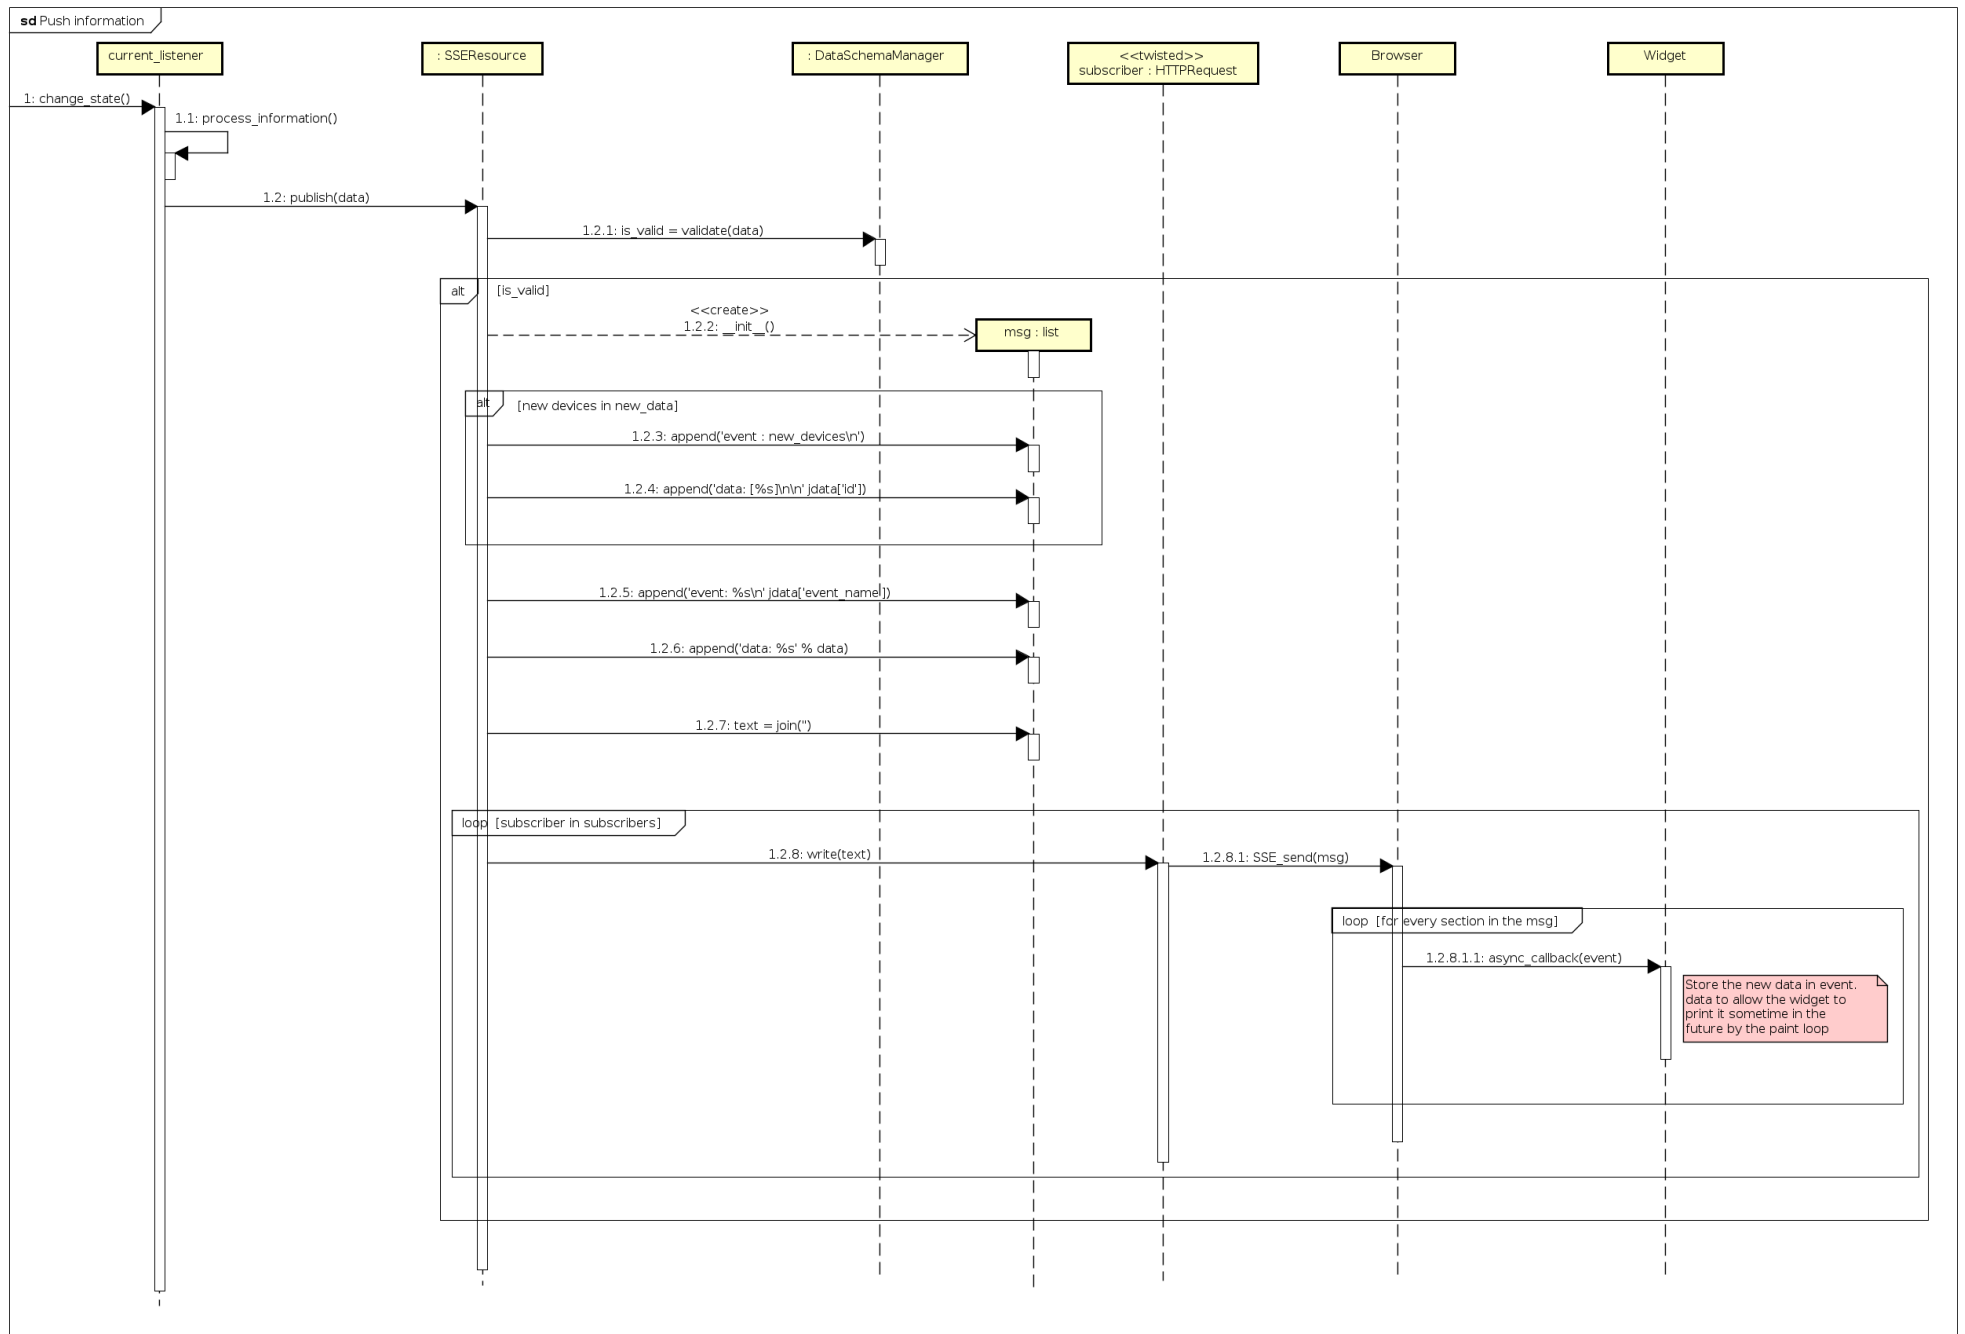


Figure 7: The backend pushes some information into the frontend

6 Methodology of work

I haven't followed strictly a methodology but a mix of them. Instead of following all the rules that a development methodology is made of, I have applied the principles it's base on to improve the quality of the software.

That way, I have paid a lot of attention to the design of the architecture before its implementation was done to ensure that the structure of the software and the interactions between its different components was clean. I have also paid a lot of attention to the testing of the software to be able to change the architecture of the software without breaking everything in the process. Finally, the documentation is also very important for the project and I have invest a lot of time writing documentation for the potential users of the framework.

Currently there are almost 100 tests in both the backend and the frontend that test every use case as a whole and in its smaller parts. In order to launch them, you only need to use trial, the default test case utility provided by twisted.

Similarly, currently there are about 30 pages of documentation that cover every aspect of the framework in a detailed way to allow programmers an easy understanding without the need to read the source code. To allow people to contribute to the project, the source code is also commented to allow its easy understanding.

I have followed the pep8 style guide in all my source code to ensure its cleanness, checking it periodically with an automated script.

During the entire development of my final project I've been using git as a version control system for my source code putting special effort on branch management of the repository to allow an easy release and build management. The project is currently uploaded to pip with every major release uploaded to pip as long as it's possible.

I have also used Redmine to report issues related to the software and also to have a record of how the project was evolving.

7 Work to do

There are several improvements to be done in the framework. The first one is to expand the number of listeners and widgets available at install time in the framework.

The documentation should also include a tutorial to build a real time web page with the basic abstractions described above.

The listeners abstraction could also be refactored to allow more than one type of listener to push information to the backend.

The frontend could also be refactored in order to allow the use of html5 canvas objects.

More granularity in the paint loop should be provided to avoid different widget types to be called at different times.

The high level languages should also be revised in order to allow more expressive and easy descriptions.

8 Conclusions

I think lupulo is a useful framework to build real time web pages that are focused on visualizing information, so I'll continue its development if there are enough active users.

I have had a lot of freedom in the development of my final project so I have decided from bottom up how to build the framework, what I should focus on, what the priorities of the development were, which methodology to use and even what my final project was altogether.

For me it has been a good experience to be able to build a project of this complexity in five months, where I had so much freedom and I could focus on development, instead that on other areas.